

# ADVANCED MACHINE LEARNING

## GAUSSIAN PROCESSES

### LECTURE 3

Mattias Villani

**Division of Statistics and Machine Learning  
Department of Computer and Information Science  
Linköping University**



# LECTURE OVERVIEW

- ▶ Lecture 3
  - ▶ Gaussian Process Classification
  - ▶ More GP models

# CLASSIFICATION WITH LOGISTIC REGRESSION

- ▶ **Classification:** binary response  $y \in \{-1, 1\}$  predicted by features  $\mathbf{x}$ .
- ▶ Example: linear logistic regression

$$Pr(y = 1|\mathbf{x}) = \lambda(\mathbf{x}^T \mathbf{w})$$

where  $\lambda(z)$  is the logistic **link function**

$$\lambda(z) = \frac{1}{1 + \exp(-z)}$$

- ▶  $\lambda(z)$  'squashes' the linear prediction  $\mathbf{x}^T \mathbf{w} \in \mathbb{R}$  into  $\lambda(\mathbf{x}^T \mathbf{w}) \in [0, 1]$ .
- ▶ Logistic regression has **linear decision boundaries**.

# GP CLASSIFICATION

- Obvious **GP extension** of logistic regression: replace  $\mathbf{x}^T \mathbf{w}$  by  $f(\mathbf{x})$  where


$$f(\mathbf{x}) \sim GP(0, k(\mathbf{x}, \mathbf{x}'))$$

and squash  $f$  through logistic function (or normal CDF)


$$Pr(y = 1|\mathbf{x}) = \lambda(f(\mathbf{x}))$$

- Decision boundaries are now non-parametric (GP). Flexible.

# SQUASHING A LINEAR FUNCTION



# SQUASHING A GP FUNCTION



# GP CLASSIFICATION - INFERENCE

- **Prediction** for a test case  $\mathbf{x}_*$ :

$$Pr(y_* = +1 | \mathbf{X}, \mathbf{y}, \mathbf{x}_*) = \int \sigma(f_*) p(f_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y}) df_*$$

where  $\sigma(f_*)$  is some sigmoidal function (logistic, normal CDF...) and  $f_*$  is the latent  $f$  at the test input  $\mathbf{x}_*$ .

- The posterior distribution of  $f_*$  is

$$p(f_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y}) = \int p(f_* | \mathbf{X}, \mathbf{x}_*, \mathbf{f}) p(\mathbf{f} | \mathbf{X}, \mathbf{y}) d\mathbf{f}$$

where

$$p(\mathbf{f} | \mathbf{X}, \mathbf{y}) \propto p(\mathbf{y} | \mathbf{f}) p(\mathbf{f} | \mathbf{X})$$

is the posterior of  $\mathbf{f}$  from the training data.

- Note that  $p(\mathbf{y} | \mathbf{f})$  is no longer Gaussian in classification problems. Posterior  $p(\mathbf{f} | \mathbf{X}, \mathbf{y})$  is not analytically tractable.

# THE LAPLACE APPROXIMATION

- ▶ Approximates  $p(\mathbf{f}|\mathbf{X}, \mathbf{y})$  with  $N(\hat{\mathbf{f}}, \mathbf{A}^{-1})$ , where  $\hat{\mathbf{f}}$  is the posterior mode and  $\mathbf{A}$  is the negative Hessian of the log posterior at  $\mathbf{f} = \hat{\mathbf{f}}$ .
- ▶ The log posterior is (proportional to)

$$\begin{aligned}\Psi(\mathbf{f}) &= \log p(\mathbf{y}|\mathbf{f}) + \log p(\mathbf{f}|\mathbf{X}) \\ &= \log p(\mathbf{y}|\mathbf{f}) - \frac{1}{2}\mathbf{f}^T K^{-1}\mathbf{f} - \frac{1}{2}\log |K| - \frac{n}{2}\log 2\pi\end{aligned}$$

- ▶ Differentiating wrt  $\mathbf{f}$

$$\begin{aligned}\nabla \Psi(\mathbf{f}) &= \nabla \log p(\mathbf{y}|\mathbf{f}) - K^{-1}\mathbf{f} \\ \nabla \nabla \Psi(\mathbf{f}) &= \nabla \nabla \log p(\mathbf{y}|\mathbf{f}) - K^{-1} = -W - K^{-1}\end{aligned}$$

where  $W$  is a diagonal matrix since each  $y_i$  only depends on its  $f_i$ .

- ▶ Use **Newton's method** to iterate to the mode.
- ▶ **Approximate predictions** of  $f_*$  are **possible**.
- ▶ Predictions of  $y_*$  require one-dimensional numerical integration.



# MARKOV CHAIN MONTE CARLO

- ▶ Metropolis-Hastings (or Hamiltonian MC) to **sample from training posterior**

$$\mathbf{f}|\mathbf{x}, \mathbf{y}, \theta$$

Produces  $\mathbf{f}^{(1)}, \dots, \mathbf{f}^{(N)}$  draws.

- ▶ For each  $\mathbf{f}^{(i)}$ , **sample the test posterior**  $\mathbf{f}_*$  from

$$\mathbf{f}_*|\mathbf{f}^{(i)}, \mathbf{x}, \mathbf{x}_* \sim N\left(K(\mathbf{x}_*, \mathbf{x})K(\mathbf{x}, \mathbf{x})^{-1}\mathbf{f}^{(i)}, K(\mathbf{x}_*, \mathbf{x}_*) - K(\mathbf{x}_*, \mathbf{x})K(\mathbf{x}, \mathbf{x})^{-1}\right)$$

Note that this does not depend on  $\mathbf{y}$  since we condition on  $\mathbf{f}$ .

Noise-free GP fit. Produces  $\mathbf{f}_*^{(1)}, \dots, \mathbf{f}_*^{(N)}$  draws.


- ▶ For each  $\mathbf{f}_*^{(i)}$ , **sample a prediction** from

$$p(\mathbf{y}_*|\mathbf{f}_*^{(i)}, \theta).$$

Produces a draws from the predictive distribution  $p(\mathbf{y}_*|\mathbf{x}_*, \mathbf{x}, \mathbf{y}, \theta)$ .


- ▶ Straightforward (at least in principle) to also **sample the hyperparameters**  $\theta$ . Slice sampling.

# IRIS DATA - SEPAL - SE KERNEL WITH ARD




Laplace:  $\hat{\ell}_1 = 1.7214, \hat{\ell}_2 = 185.5040, \sigma_f = 1.4361$

EP:  $\hat{\ell}_1 = 1.7189, \hat{\ell}_2 = 55.5003, \sigma_f = 1.4343$




# IRIS DATA - PETAL - SE KERNEL WITH ARD




Laplace:  $\hat{\ell}_1 = 1.7606, \hat{\ell}_2 = 0.8804, \sigma_f = 4.9129$

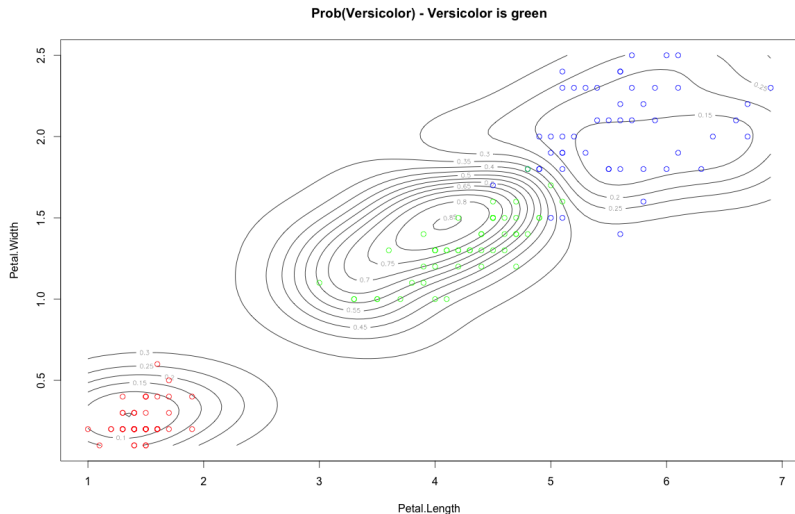
EP:  $\hat{\ell}_1 = 2.1139, \hat{\ell}_2 = 1.0720, \sigma_f = 5.3369$



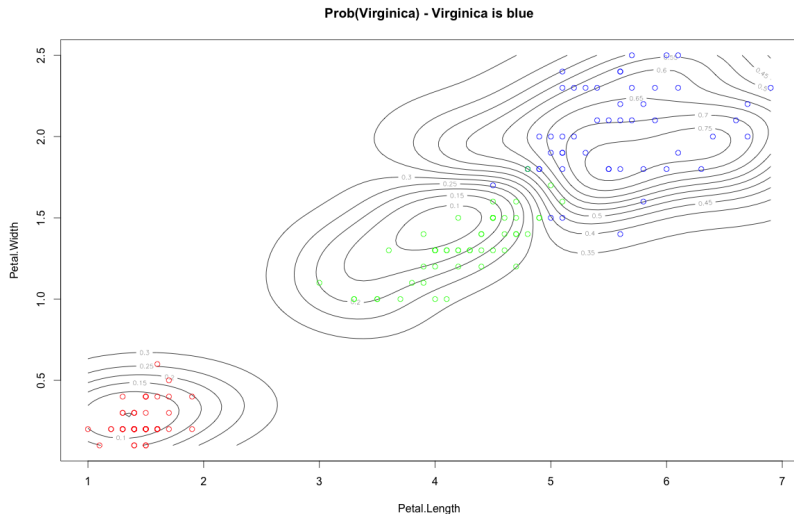
# IRIS DATA - PETAL - ALL THREE CLASSES




# IRIS DATA - PETAL - ALL THREE CLASSES




# IRIS DATA - PETAL - ALL THREE CLASSES




# IRIS DATA - PETAL - DECISION BOUNDARIES



# TOY DATA 1 - SE KERNEL WITH ARD




$$\text{EP: } \hat{\ell}_1 = 2.4503, \hat{\ell}_2 = 721.7405, \sigma_f = 4.7540$$







## TOY DATA 2 - SE KERNEL WITH ARD



EP:  $\hat{\ell}_1 = 8.3831, \hat{\ell}_2 = 1.9587, \sigma_f = 4.5483$




# TOY DATA 3 - SE KERNEL WITH ARD



Laplace:  $\hat{\ell}_1 = 0.7726, \hat{\ell}_2 = 0.6974, \sigma_f = 11.7854$

EP:  $\hat{\ell}_1 = 1.2685, \hat{\ell}_2 = 1.0941, \sigma_f = 17.2774$



# GAUSSIAN PROCESS OPTIMIZATION (GPO)

- ▶ **Aim:** minimization of **expensive** function

$$\operatorname{argmin}_{\mathbf{x} \in \mathcal{X}} f(\mathbf{x})$$


- ▶ Typical applications: **hyperparameter estimation**.

- ▶ **GPO idea:**

- ▶ Assign GP prior to the unknown function  $f$ .
- ▶ Evaluate the function at some values  $x_1, x_2, \dots, x_n$ .
- ▶ Update to posterior  $f|x_1, \dots, x_n \sim GP(\mu, K)$ . Noise-free model.
- ▶ Use the GP posterior of  $f$  to find a new evaluation point  $x_{n+1}$ .  
**Explore** vs **Exploit**.
- ▶ Iterate until the change in optimum is lower than some tolerance.

- ▶ **Bayesian Optimization**. Bayesian Numerics. Probabilistic numerics.

# EXPLORE-EXPLOIT ILLUSTRATION



# ACQUISITION FUNCTIONS

## ► Probability of Improvement (PI)

$$a_{PI}(\mathbf{x}; \{\mathbf{x}_n, y_n\}, \theta) \equiv \Pr(f(\mathbf{x}) < f(\mathbf{x}_{best})) = \Phi(\gamma(\mathbf{x}))$$

where

$$\gamma(\mathbf{x}) = \frac{f(\mathbf{x}_{best}) - \mu(\mathbf{x}; \{\mathbf{x}_n, y_n\}, \theta)}{\sigma(\mathbf{x}; \{\mathbf{x}_n, y_n\}, \theta)}$$

## ► Expected Improvement (EI)


$$a_{EI}(\mathbf{x}; \{\mathbf{x}_n, y_n\}, \theta) = \sigma(\mathbf{x}; \{\mathbf{x}_n, y_n\}, \theta) [\gamma(\mathbf{x})\Phi(\gamma(\mathbf{x})) + \mathcal{N}(\gamma(\mathbf{x}); 0, 1)]$$

## ► Lower Confidence Bound (LCB)


$$a_{EI}(\mathbf{x}; \{\mathbf{x}_n, y_n\}, \theta) = \mu(\mathbf{x}; \{\mathbf{x}_n, y_n\}, \theta) - \kappa \cdot \sigma(\mathbf{x}; \{\mathbf{x}_n, y_n\}, \theta)$$

- Note: need to maximize the acquisition function to choose  $\mathbf{x}_{next}$ .  
Non-convex, but cheaper and simpler than original problem.

# ACQUISITION FUNCTIONS FROM BROCHU ET AL



# CONVNETS - SNOEK ET AL (NIPS, 2012)



# MORE GP MODELS

## ► Heteroscedastic GP regression

$$y = f(x) + \exp[g(x)] \epsilon$$

so where  $f \sim GP[m_f(x), k_f(x, x')]$  independently of

$$g \sim GP[m_g(x), k_g(x, x')].$$

## ► GP for density estimation

$$p(x) = \frac{\exp[f(x)]}{\int_{\mathbb{R}} \exp[f(t)] dt}$$

where  $f \sim GP[m(x), k(x, x')]$ . Appealing mean function:

$$m(x) = -\frac{1}{2\theta_2}(x - \theta_1)^2 \text{ [i.e. best guess is a normal density].}$$

## ► Shared latent GP for dependent multivariate data ( $k \ll p$ )

$$\begin{pmatrix} y_1(x) \\ \vdots \\ y_p(x) \end{pmatrix} = \mathbf{L}_{p \times k} \begin{pmatrix} f_1(x) \\ \vdots \\ f_k(x) \end{pmatrix} + \begin{pmatrix} g_1(x) \\ \vdots \\ g_p(x) \end{pmatrix} + \begin{pmatrix} \epsilon_1 \\ \vdots \\ \epsilon_p \end{pmatrix}$$