

---

# PROBABILISTIC ROBOTICS

---

**Sebastian THRUN**

*Stanford University*

*Stanford, CA*



**Wolfram BURGARD**

*University of Freiburg*

*Freiburg, Germany*



**Dieter FOX**

*University of Washington*

*Seattle, WA*

EARLY DRAFT—NOT FOR DISTRIBUTION  
©Sebastian Thrun, Dieter Fox, Wolfram Burgard, 1999-2000

---

# CONTENTS

<b>1</b>	<b>INTRODUCTION</b>	<b>1</b>
1.1	Uncertainty in Robotics	1
1.2	Probabilistic Robotics	3
1.3	Implications	5
1.4	Road Map	6
1.5	Bibliographical Remarks	7
<b>2</b>	<b>RECURSIVE STATE ESTIMATION</b>	<b>9</b>
2.1	Introduction	9
2.2	Basic Concepts in Probability	10
2.3	Robot Environment Interaction	16
2.3.1	State	16
2.3.2	Environment Interaction	18
2.3.3	Probabilistic Generative Laws	20
2.3.4	Belief Distributions	22
2.4	Bayes Filters	23
2.4.1	The Bayes Filter Algorithm	23
2.4.2	Example	24
2.4.3	Mathematical Derivation of the Bayes Filter	28
2.4.4	The Markov Assumption	30
2.5	Representation and Computation	30
2.6	Summary	31
2.7	Bibliographical Remarks	32
<b>3</b>	<b>GAUSSIAN FILTERS</b>	<b>33</b>
3.1	Introduction	33
3.2	The Kalman Filter	34

3.2.1	Linear Gaussian Systems	34
3.2.2	The Kalman Filter Algorithm	36
3.2.3	Illustration	37
3.2.4	Mathematical Derivation of the KF	39
3.3	The Extended Kalman Filter	48
3.3.1	Linearization Via Taylor Expansion	49
3.3.2	The EKF Algorithm	50
3.3.3	Mathematical Derivation of the EKF	51
3.3.4	Practical Considerations	53
3.4	The Information Filter	55
3.4.1	Canonical Representation	55
3.4.2	The Information Filter Algorithm	57
3.4.3	Mathematical Derivation of the Information Filter	58
3.4.4	The Extended Information Filter Algorithm	60
3.4.5	Mathematical Derivation of the Extended Information Filter	61
3.4.6	Practical Considerations	62
3.5	Summary	64
3.6	Bibliographical Remarks	65
<b>4</b>	<b>NONPARAMETRIC FILTERS</b>	<b>67</b>
4.1	The Histogram Filter	68
4.1.1	The Discrete Bayes Filter Algorithm	69
4.1.2	Continuous State	69
4.1.3	Decomposition Techniques	73
4.1.4	Binary Bayes Filters With Static State	74
4.2	The Particle Filter	77
4.2.1	Basic Algorithm	77
4.2.2	Importance Sampling	80
4.2.3	Mathematical Derivation of the PF	82
4.2.4	Properties of the Particle Filter	84
4.3	Summary	89
4.4	Bibliographical Remarks	90
<b>5</b>	<b>ROBOT MOTION</b>	<b>91</b>
5.1	Introduction	91

### 2.3.3 Probabilistic Generative Laws

The evolution of state and measurements is governed by probabilistic laws. In general, the state at time  $x_t$  is generated stochastically. Thus, it makes sense to specify the probability distribution from which  $x_t$  is generated. At first glance, the emergence of state  $x_t$  might be conditioned on all past states, measurements, and controls. Hence, the probabilistic law characterizing the evolution of state might be given by a probability distribution of the following form:

$$p(x_t \mid x_{0:t-1}, z_{1:t-1}, u_{1:t}) \quad (2.32)$$

(Notice that through no particular motivation we assume here that the robot executes a control action  $u_1$  first, and then takes a measurement  $z_1$ .) However, if the state  $x$  is

complete then it is a sufficient summary of all that happened in previous time steps. In particular,  $x_{t-1}$  is a sufficient statistic of all previous controls and measurements up to this point, that is,  $u_{1:t-1}$  and  $z_{1:t-1}$ . From all the variables in the expression above, only the control  $u_t$  matters if we know the state  $x_{t-1}$ . In probabilistic terms, this insight is expressed by the following equality:

$$p(x_t \mid x_{0:t-1}, z_{1:t-1}, u_{1:t}) = p(x_t \mid x_{t-1}, u_t) \quad (2.33)$$

The property expressed by this equality is an example of *conditional independence*. It states that certain variables are independent of others if one knows the values of a third group of variables, the conditioning variables. Conditional independence will be exploited pervasively in this book, as it is the main source of tractability of probabilistic robotics algorithms.

Similarly, one might want to model the process by which measurements are being generated. Again, if  $x_t$  is complete, we have an important conditional independence:

$$p(z_t \mid x_{0:t}, z_{1:t-1}, u_{1:t}) = p(z_t \mid x_t) \quad (2.34)$$

In other words, the state  $x_t$  is sufficient to predict the (potentially noisy) measurement  $z_t$ . Knowledge of any other variable, such as past measurements, controls or even past states, is irrelevant if  $x_t$  is complete.

This discussion leaves open as to what the two resulting conditional probabilities are:  $p(x_t \mid x_{t-1}, u_t)$  and  $p(z_t \mid x_t)$ . The probability  $p(x_t \mid x_{t-1}, u_t)$  is the *state transition probability*. It specifies how environmental state evolves over time as a function of robot controls  $u_t$ . Robot environments are stochastic, which is reflected by the fact that  $p(x_t \mid x_{t-1}, u_t)$  is a probability distribution, not a deterministic function. Sometimes the state transition distribution does not depend on the time index  $t$ , in which case we may write it as  $p(x' \mid u, x)$ , where  $x'$  is the successor and  $x$  the predecessor state.

The probability  $p(z_t \mid x_t)$  is called the *measurement probability*. It also may not depend on the time index  $t$ , in which case it shall be written as  $p(z \mid x)$ . The measurement probability specifies the probabilistic law according to which measurements  $z$  are generated from the environment state  $x$ . Measurements are usually noisy projections of the state.

The state transition probability and the measurement probability together describe the dynamical stochastic system of the robot and its environment. Figure ?? illustrates the evolution of states and measurements, defined through those probabilities. The

state at time  $t$  is stochastically dependent on the state at time  $t - 1$  and the control  $u_t$ . The measurement  $z_t$  depends stochastically on the state at time  $t$ . Such a temporal generative model is also known as hidden Markov model (HMM) or dynamic Bayes network (DBN). To specify the model fully, we also need an initial state distribution  $p(x_0)$ .

### 2.3.4 Belief Distributions

Another key concept in probabilistic robotics is that of a *belief*. A belief reflects the robot's internal knowledge about the state of the environment. We already discussed that state cannot be measured directly. For example, a robot's pose might be  $x = \langle 14.12, 12.7, 0.755 \rangle$  in some global coordinate system, but it usually cannot know its pose, since poses are not measurable directly (not even with GPS!). Instead, the robot must infer its pose from data. We therefore distinguish the true state from its internal *belief*, or *state of knowledge* with regards to that state.

Probabilistic robotics represents beliefs through conditional probability distributions. A belief distribution assigns a probability (or density value) to each possible hypothesis with regards to the true state. Belief distributions are posterior probabilities over state variables conditioned on the available data. We will denote belief over a state variable  $x_t$  by  $bel(x_t)$ , which is an abbreviation for the posterior

$$bel(x_t) = p(x_t \mid z_{1:t}, u_{1:t}) . \quad (2.35)$$

This posterior is the probability distribution over the state  $x_t$  at time  $t$ , conditioned on all past measurements  $z_{1:t}$  and all past controls  $u_{1:t}$ .

The reader may notice that we silently assume that the belief is taken *after* incorporating the measurement  $z_t$ . Occasionally, it will prove useful to calculate a posterior *before* incorporating  $z_t$ , just after executing the control  $u_t$ . Such a posterior will be denoted as follows:

$$\overline{bel}(x_t) = p(x_t \mid z_{1:t-1}, u_{1:t}) \quad (2.36)$$

This probability distribution is often referred to as *prediction* in the context of probabilistic filtering. This terminology reflects the fact that  $\overline{bel}(x_t)$  predicts the state at time  $t$  based on the previous state posterior, before incorporating the measurement at time  $t$ . Calculating  $bel(x_t)$  from  $\overline{bel}(x_t)$  is called *correction* or the *measurement update*.

## 2.4 BAYES FILTERS

### 2.4.1 The Bayes Filter Algorithm

The most general algorithm for calculating beliefs is given by the *Bayes filter* algorithm. This algorithm calculates the belief distribution  $bel$  from measurement and control data. We will first state the basic algorithm and elucidate it with a numerical example. After that, we will derive it mathematically from the assumptions made so far.

Table 2.1 depicts the basic Bayes filter in pseudo-algorithmic form. The Bayes filter is recursive, that is, the belief  $bel(x_t)$  at time  $t$  is calculated from the belief  $bel(x_{t-1})$  at time  $t-1$ . Its input is the belief  $bel$  at time  $t-1$ , along with the most recent control  $u_t$  and the most recent measurement  $z_t$ . Its output is the belief  $bel(x_t)$  at time  $t$ . Table 2.1 only depicts a single step of the Bayes Filter algorithm: the *update rule*. This update rule is applied recursively, to calculate the belief  $bel(x_t)$  from the belief  $bel(x_{t-1})$ , calculated previously.

The Bayes filter algorithm possesses two essential steps. In Line 3, it processes the control  $u_t$ . It does so by calculating a belief over the state  $x_t$  based on the prior belief over state  $x_{t-1}$  and the control  $u_t$ . In particular, the belief  $\overline{bel}(x_t)$  that the robot assigns to state  $x_t$  is obtained by the integral (sum) of the product of two distributions: the prior assigned to  $x_{t-1}$ , and the probability that control  $u_t$  induces a transition from  $x_{t-1}$  to  $x_t$ . The reader may recognize the similarity of this update step to Equation (2.12). As noted above, this update step is called the control update, or *prediction*.

The second step of the Bayes filter is called the measurement update. In Line 4, the Bayes filter algorithm multiplies the belief  $\overline{bel}(x_t)$  by the probability that the measurement  $z_t$  may have been observed. It does so for each hypothetical posterior state  $x_t$ . As will become apparent further below when actually deriving the basic filter equations, the resulting product is generally not a probability, that is, it may not integrate to 1. Hence, the result is normalized, by virtue of the normalization constant  $\eta$ . This leads to the final belief  $bel(x_t)$ , which is returned in Line 6 of the algorithm.

To compute the posterior belief recursively, the algorithm requires an initial belief  $bel(x_0)$  at time  $t = 0$  as boundary condition. If one knows the value of  $x_0$  with certainty,  $bel(x_0)$  should be initialized with a point mass distribution that centers all probability mass on the correct value of  $x_0$ , and assigns zero probability anywhere else. If one is entirely ignorant about the initial value  $x_0$ ,  $bel(x_0)$  may be initialized using a uniform distribution over the domain of  $x_0$  (or related distribution from the Dirichlet family of distributions). Partial knowledge of the initial value  $x_0$  can be

```

1:   Algorithm Bayes_filter( $bel(x_{t-1}), u_t, z_t$ ):
2:     for all  $x_t$  do
3:        $\overline{bel}(x_t) = \int p(x_t \mid u_t, x_{t-1}) bel(x_{t-1}) dx$ 
4:        $bel(x_t) = \eta p(z_t \mid x_t) \overline{bel}(x_t)$ 
5:     endfor
6:     return  $bel(x_t)$ 

```

**Table 2.1** The general algorithm for Bayes filtering.

expressed by non-uniform distributions; however, the two cases of full knowledge and full ignorance are the most common ones in practice.

The algorithm Bayes filter can only be implemented in the form stated here for very simple estimation problems. In particular, we either need to be able to carry out the integration in Line 3 and the multiplication in Line 4 in closed form, or we need to restrict ourselves to finite state spaces, so that the integral in Line 3 becomes a (finite) sum.



### 2.4.3 Mathematical Derivation of the Bayes Filter

The correctness of the Bayes filter algorithm is shown by induction. To do so, we need to show that it correctly calculates the posterior distribution  $p(x_t \mid z_{1:t}, u_{1:t})$  from the corresponding posterior one time step earlier,  $p(x_{t-1} \mid z_{1:t-1}, u_{1:t-1})$ . The correctness follows then by induction under the assumption that we correctly initialized the prior belief  $bel(x_0)$  at time  $t = 0$ .

Our derivation requires that the state  $x_t$  is complete, as defined in Section 2.3.1, and it requires that controls are chosen at random. The first step of our derivation involves the application of Bayes rule (2.23) to the target posterior:

$$\begin{aligned} p(x_t \mid z_{1:t}, u_{1:t}) &= \frac{p(z_t \mid x_t, z_{1:t-1}, u_{1:t}) p(x_t \mid z_{1:t-1}, u_{1:t})}{p(z_t \mid z_{1:t-1}, u_{1:t})} \\ &= \eta p(z_t \mid x_t, z_{1:t-1}, u_{1:t}) p(x_t \mid z_{1:t-1}, u_{1:t}) \end{aligned} \quad (2.55)$$

We now exploit the assumption that our state is complete. In Section 2.3.1, we defined a state  $x_t$  to be complete if no variables prior to  $x_t$  may influence the stochastic evolution of future states. In particular, if we (hypothetically) knew the state  $x_t$  and were interested in predicting the measurement  $z_t$ , no past measurement or control would provide us additional information. In mathematical terms, this is expressed by the following conditional independence:

$$p(z_t \mid x_t, z_{1:t-1}, u_{1:t}) = p(z_t \mid x_t). \quad (2.56)$$

Such a statement is another example of *conditional independence*. It allows us to simplify (2.55) as follows:

$$p(x_t \mid z_{1:t}, u_{1:t}) = \eta p(z_t \mid x_t) p(x_t \mid z_{1:t-1}, u_{1:t}) \quad (2.57)$$

and hence

$$bel(x_t) = \eta p(z_t | x_t) \overline{bel}(x_t) \quad (2.58)$$

This equation is implemented in Line 4 of the Bayes filter algorithm in Table 2.1.

Next, we expand the term  $\overline{bel}(x_t)$ , using (2.12):

$$\begin{aligned} \overline{bel}(x_t) &= p(x_t | z_{1:t-1}, u_{1:t}) \\ &= \int p(x_t | x_{t-1}, z_{1:t-1}, u_{1:t}) p(x_{t-1} | z_{1:t-1}, u_{1:t}) dx_{t-1} \end{aligned} \quad (2.59)$$

Once again, we exploit the assumption that our state is complete. This implies if we know  $x_{t-1}$ , past measurements and controls convey no information regarding the state  $x_t$ . This gives us

$$p(x_t | x_{t-1}, z_{1:t-1}, u_{1:t}) = p(x_t | x_{t-1}, u_t) \quad (2.60)$$

Here we retain the control variable  $u_t$ , since it does *not* predate the state  $x_{t-1}$ . Finally, we note that the control  $u_t$  can safely be omitted from the set of conditioning variables in  $p(x_{t-1} | z_{1:t-1}, u_{1:t})$  for randomly chosen controls. This gives us the recursive update equation

$$\overline{bel}(x_t) = \int p(x_t | x_{t-1}, u_t) p(x_{t-1} | z_{1:t-1}, u_{1:t-1}) dx_{t-1} \quad (2.61)$$

As the reader easily verifies, this equation is implemented by Line 3 of the Bayes filter algorithm in Table 2.1. To summarize, the Bayes filter algorithm calculates the posterior over the state  $x_t$  conditioned on the measurement and control data up to time  $t$ . The derivation assumes that the world is Markov, that is, the state is complete.

Any concrete implementation of this algorithm requires three probability distributions: The initial belief  $p(x_0)$ , the measurement probability  $p(z_t | x_t)$ , and the state transition probability  $p(x_t | u_t, x_{t-1})$ . We have not yet specified these densities, but will do so in later chapters (Chapters 5 and ??). Additionally, we also need a representation for the belief  $bel(x_t)$ , which will also be discussed further below.

# 3

---

## GAUSSIAN FILTERS

### 3.1 INTRODUCTION

This chapter describes an important family of recursive state estimators, collectively called *Gaussian Filters*. Historically, Gaussian filters constitute the earliest tractable implementations of the Bayes filter for continuous spaces. They are also by far the most popular family of techniques to date—despite a number of shortcomings.

Gaussian techniques all share the basic idea that beliefs are represented by multivariate normal distributions. We already encountered a definition of the multivariate normal distribution in Equation (2.4), which is restated here:

$$p(x) = \det(2\pi\Sigma)^{-\frac{1}{2}} \exp\left\{-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right\} \quad (3.1)$$

This density over the variable  $x$  is characterized by two sets of parameters: The mean  $\mu$  and the covariance  $\Sigma$ . The mean  $\mu$  is a vector that possesses the same dimensionality as the state  $x$ . The covariance is a quadratic matrix that is symmetric and positive-semidefinite. Its dimension is the dimensionality of the state  $x$  squared. Thus, the number of elements in the covariance matrix depends quadratically on the number of elements in the state vector.

The commitment to represent the posterior by a Gaussian has important ramifications. Most importantly, Gaussians are unimodal, that is, they possess a single maximum. Such a posterior is characteristic of many tracking problems in robotics, in which the posterior is focused around the true state with a small margin of uncertainty. Gaussian posteriors are a poor match for many global estimation problems in which many distinct hypotheses exist, each of which forming its own mode in the posterior.

The representation of a Gaussian by its mean and covariance is called the *moments representation*. This is because the mean and covariance are the first and second moments of a probability distribution; all other moments are zero for normal distributions. In this chapter, we will also discuss an alternative representation, called *canonical representation*, or sometimes *natural representation*. Both representations, the moments and the canonical representations, are functionally equivalent in that a bijective mapping exists that transforms one into the other (and back). However, they lead to filter algorithms with orthogonal computational characteristics.

This chapter introduces the two basic Gaussian filter algorithms.

- Section 3.2 describes the Kalman filter, which implements the Bayes filter using the moments representation for a restricted class of problems with linear dynamics and measurement functions.
- The Kalman filter is extended to nonlinear problems in Section 3.3, which describes the extended Kalman filter.
- Section 3.4 describes the information filter, which is the dual of the Kalman filter using the canonical representation of Gaussians.

## 3.2 THE KALMAN FILTER

### 3.2.1 Linear Gaussian Systems

Probably the best studied technique for implementing Bayes filters is the *Kalman filter* (KF). The Kalman filter was invented in the 1950s by Rudolph Emil Kalman, as a technique for filtering and prediction in linear systems. The Kalman filter implements belief computation for continuous states. It is not applicable to discrete or hybrid state spaces.

The Kalman filter represents beliefs by the moments representation: At time  $t$ , the belief is represented by the mean  $\mu_t$  and the covariance  $\Sigma_t$ . Posteriors are Gaussian if the following three properties hold, in addition to the Markov assumptions of the Bayes filter.

1. The next state probability  $p(x_t \mid u_t, x_{t-1})$  must be a *linear* function in its arguments with added Gaussian noise. This is expressed by the following equation:

$$x_t = A_t x_{t-1} + B_t u_t + \varepsilon_t . \quad (3.2)$$

Here  $x_t$  and  $x_{t-1}$  are state vectors, and  $u_t$  is the control vector at time  $t$ . In our notation, both of these vectors are vertical vectors, that is, they are of the form

$$x_t = \begin{pmatrix} x_{1,t} \\ x_{2,t} \\ \vdots \\ x_{n,t} \end{pmatrix} \quad \text{and} \quad u_t = \begin{pmatrix} u_{1,t} \\ u_{2,t} \\ \vdots \\ u_{m,t} \end{pmatrix}. \quad (3.3)$$

$A_t$  and  $B_t$  are matrices.  $A_t$  is a square matrix of size  $n \times n$ , where  $n$  is the dimension of the state vector  $x_t$ .  $B_t$  is of size  $n \times m$ , with  $m$  being the dimension of the control vector  $u_t$ . By multiplying the state and control vector with the matrices  $A_t$  and  $B_t$ , respectively, the state transition function becomes *linear* in its arguments. Thus, Kalman filters assume linear system dynamics.

The random variable  $\varepsilon_t$  in (3.2) is a Gaussian random vector that models the randomness in the state transition. It is of the same dimension as the state vector. Its mean is zero and its covariance will be denoted  $R_t$ . A state transition probability of the form (3.2) is called a *linear Gaussian*, to reflect the fact that it is linear in its arguments with additive Gaussian noise.

Equation (3.2) defines the state transition probability  $p(x_t | u_t, x_{t-1})$ . This probability is obtained by plugging Equation (3.2) into the definition of the multivariate normal distribution (3.1). The mean of the posterior state is given by  $A_t x_{t-1} + B_t u_t$  and the covariance by  $R_t$ :

$$\begin{aligned} p(x_t | u_t, x_{t-1}) &= \det(2\pi R_t)^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (x_t - A_t x_{t-1} - B_t u_t)^T R_t^{-1} (x_t - A_t x_{t-1} - B_t u_t) \right\} \end{aligned} \quad (3.4)$$

2. The measurement probability  $p(z_t | x_t)$  must also be *linear* in its arguments, with added Gaussian noise:

$$z_t = C_t x_t + \delta_t. \quad (3.5)$$

Here  $C_t$  is a matrix of size  $k \times n$ , where  $k$  is the dimension of the measurement vector  $z_t$ . The vector  $\delta_t$  describes the measurement noise. The distribution of  $\delta_t$  is a multivariate Gaussian with zero mean and covariance  $Q_t$ . The measurement probability is thus given by the following multivariate normal distribution:

$$p(z_t | x_t) = \det(2\pi Q_t)^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (z_t - C_t x_t)^T Q_t^{-1} (z_t - C_t x_t) \right\} \quad (3.6)$$

```

1:  Algorithm Kalman_filter( $\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$ ):
2:       $\bar{\mu}_t = A_t \mu_{t-1} + B_t u_t$ 
3:       $\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t$ 
4:       $K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1}$ 
5:       $\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$ 
6:       $\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$ 
7:      return  $\mu_t, \Sigma_t$ 

```

**Table 3.1** The Kalman filter algorithm for linear Gaussian state transitions and measurements.

3. Finally, the initial belief  $bel(x_0)$  must be normal distributed. We will denote the mean of this belief by  $\mu_0$  and the covariance by  $\Sigma_0$ :

$$bel(x_0) = p(x_0) = \det(2\pi\Sigma_0)^{-\frac{1}{2}} \exp\left\{-\frac{1}{2}(x_0 - \mu_0)^T \Sigma_0^{-1}(x_0 - \mu_0)\right\}$$

These three assumptions are sufficient to ensure that the posterior  $bel(x_t)$  is always a Gaussian, for any point in time  $t$ . The proof of this non-trivial result can be found below, in the mathematical derivation of the Kalman filter (Section 3.2.4).

### 3.2.2 The Kalman Filter Algorithm

The Kalman filter algorithm is depicted in Table 3.1. Kalman filters represent the belief  $bel(x_t)$  at time  $t$  by the mean  $\mu_t$  and the covariance  $\Sigma_t$ . The input of the Kalman filter is the belief at time  $t - 1$ , represented by  $\mu_{t-1}$  and  $\Sigma_{t-1}$ . To update these parameters, Kalman filters require the control  $u_t$  and the measurement  $z_t$ . The output is the belief at time  $t$ , represented by  $\mu_t$  and  $\Sigma_t$ .

In Lines 2 and 3, the predicted belief  $\bar{\mu}$  and  $\bar{\Sigma}$  is calculated representing the belief  $\bar{bel}(x_t)$  one time step later, but before incorporating the measurement  $z_t$ . This belief is obtained by incorporating the control  $u_t$ . The mean is updated using the deterministic version of the state transition function (3.2), with the mean  $\mu_{t-1}$  substituted for the state  $x_{t-1}$ . The update of the covariance considers the fact that states depend on previous states through the linear matrix  $A_t$ . This matrix is multiplied twice into the covariance, since the covariance is a quadratic matrix.


The belief  $\overline{bel}(x_t)$  is subsequently transformed into the desired belief  $bel(x_t)$  in Lines 4 through 6, by incorporating the measurement  $z_t$ . The variable  $K_t$ , computed in Line 4 is called *Kalman gain*. It specifies the degree to which the measurement is incorporated into the new state estimate. Line 5 manipulates the mean, by adjusting it in proportion to the Kalman gain  $K_t$  and the deviation of the actual measurement,  $z_t$ , and the measurement predicted according to the measurement probability (3.5). Finally, the new covariance of the posterior belief is calculated in Line 6, adjusting for the information gain resulting from the measurement.

The Kalman filter is computationally quite efficient. For today's best algorithms, the complexity of matrix inversion is approximately  $O(d^{2.8})$  for a matrix of size  $d \times d$ . Each iteration of the Kalman filter algorithm, as stated here, is lower bounded by (approximately)  $O(k^{2.8})$ , where  $k$  is the dimension of the measurement vector  $z_t$ . This (approximate) cubic complexity stems from the matrix inversion in Line 4. It is also at least in  $O(n^2)$ , where  $n$  is the dimension of the state space, due to the multiplication in Line 6 (the matrix  $K_t C_t$  may be sparse). In many applications—such as the robot mapping applications discussed in later chapters—the measurement space is much lower dimensional than the state space, and the update is dominated by the  $O(n^2)$  operations.

### 3.2.3 Illustration

Figure 3.2 illustrates the Kalman filter algorithm for a simplistic one-dimensional localization scenario. Suppose the robot moves along the horizontal axis in each diagram in Figure 3.2. Let the prior over the robot location be given by the normal distribution shown in Figure 3.2a. The robot queries its sensors on its location (e.g., a GPS system), and those return a measurement that is centered at the peak of the bold Gaussian in Figure 3.2b. This bold Gaussian illustrates this measurement: Its peak is the value predicted by the sensors, and its width (variance) corresponds to the uncertainty in the measurement. Combining the prior with the measurement, via Lines 4 through 6 of the Kalman filter algorithm in Table 3.1, yields the bold Gaussian in Figure 3.2c. This belief's mean lies between the two original means, and its uncertainty radius is smaller than both contributing Gaussians. The fact that the residual uncertainty is smaller than the contributing Gaussians may appear counter-intuitive, but it is a general characteristic of information integration in Kalman filters.

Next, assume the robot moves towards the right. Its uncertainty grows due to the fact that the next state transition is stochastic. Lines 2 and 3 of the Kalman filter provides us with the Gaussian shown in bold in Figure 3.2d. This Gaussian is shifted by the amount the robot moved, and it is also wider for the reasons just explained. Next, the



**Figure 3.2** Illustration of Kalman filters: (a) initial belief, (b) a measurement (in bold) with the associated uncertainty, (c) belief after integrating the measurement into the belief using the Kalman filter algorithm, (d) belief after motion to the right (which introduces uncertainty), (e) a new measurement with associated uncertainty, and (f) the resulting belief.



robot receives a second measurement illustrated by the bold Gaussian in Figure 3.2e, which leads to the posterior shown in bold in Figure 3.2f.

As this example illustrates, the Kalman filter alternates a *measurement update step* (Lines 5-7), in which sensor data is integrated into the present belief, with a *prediction step* (or control update step), which modifies the belief in accordance to an action. The update step decreases and the prediction step increases uncertainty in the robot's belief.